

# Group10 Second Assignment

December 15, 2025

## Authors

Mattia Zanin - [mattia.zanin@studenti.unipd.it](mailto:mattia.zanin@studenti.unipd.it)  
Matteo Giorgi - [matteo.giorgi.1@studenti.unipd.it](mailto:matteo.giorgi.1@studenti.unipd.it)  
Enrico Zanello - [enrico.zanello@studenti.unipd.it](mailto:enrico.zanello@studenti.unipd.it)  
Luca Lo Buono - [luca.lobuono@studenti.unipd.it](mailto:luca.lobuono@studenti.unipd.it)

## 0.1 Plot of the time series

The first step of the analysis consists of plotting the raw time series in order to gain an initial understanding of its main characteristics.

From the plot, the series appears to fluctuate around a roughly constant level, without an evident long-term upward or downward trend. The variability of the observations seems relatively stable over time, with no clear signs of volatility clustering or explosive behavior.

Moreover, the series exhibits a recurrent oscillatory pattern, suggesting the possible presence of a seasonal or cyclical component. This visual inspection does not provide definitive evidence of stationarity, but it does not rule it out either.

At this stage, the plot serves as a preliminary diagnostic tool. A formal assessment of stationarity will be carried out in the next step using both visual techniques and statistical tests.

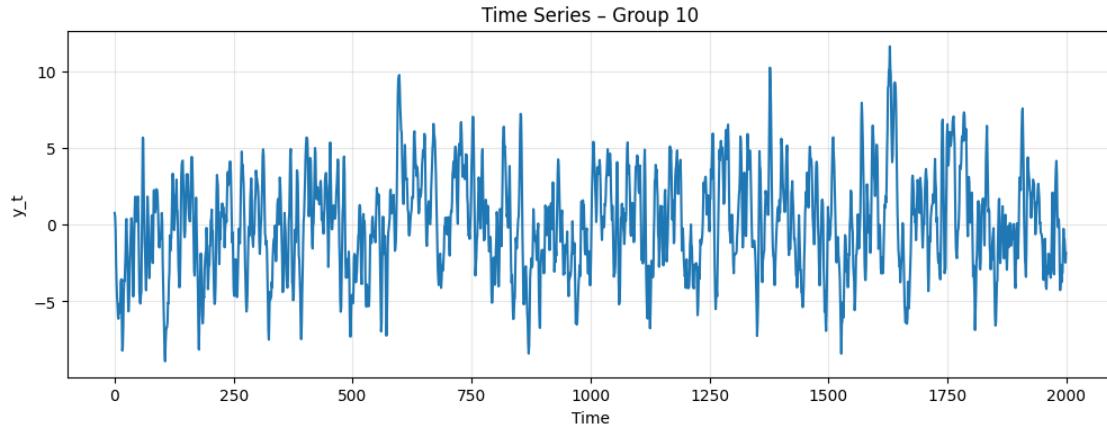
```
[1]: import pandas as pd
import matplotlib.pyplot as plt

# Load the new dataset
df = pd.read_csv("arma12_seasonal_controlled_15series.csv")

# Select Group 10
y = df["series_10"].astype(float)

# Plot the series
plt.figure(figsize=(12, 4))
plt.plot(y)
plt.title("Time Series - Group 10")
plt.xlabel("Time")
plt.ylabel("y_t")
plt.grid(alpha=0.3)
```

```
plt.show()
```



## 0.2 Check for stationarity

```
[2]: from statsmodels.tsa.stattools import adfuller

# Perform the Augmented Dickey-Fuller test
# The adfuller function returns a tuple of statistics
result = adfuller(y, 21)

# Extract and print the results
adf_stat = result[0]
p_value = result[1]
usedlag = result[2]
nobs = result[3]
critical_values = result[4]

print(f"ADF Statistic: {adf_stat}")
print(f"p-value: {p_value}")
print("Critical Values:")
for key, value in critical_values.items():
    print(f"    {key}: {value}")

# Check stationarity
if p_value < 0.05:
    print("Result: The series is Stationary")
else:
    print("Result: The series is Non-Stationary")
```

ADF Statistic: -9.783542032993036

p-value: 6.635302025669602e-17

Critical Values:

```
1%: -3.433643643742798
5%: -2.862994949652858
10%: -2.5675445538118042
Result: The series is Stationary
```

The stationarity of the series was formally assessed using the Augmented Dickey–Fuller (ADF) test.

The null hypothesis of the ADF test is that the series contains a unit root, i.e. it is non-stationary. The test returns a very small p-value, well below standard significance levels. Therefore, we reject the null hypothesis and conclude that the series is stationary.

As a consequence, no transformation (such as differencing) is required in order to achieve stationarity, and the series can be directly modeled using ARMA processes.

This result is consistent with the concept of **weak stationarity**, which requires the mean and variance of the process to be finite and constant over time, and the autocovariance structure to depend only on the lag and not on time itself. No assumptions are made on higher-order moments or on the equality of full joint distributions.

```
[3]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

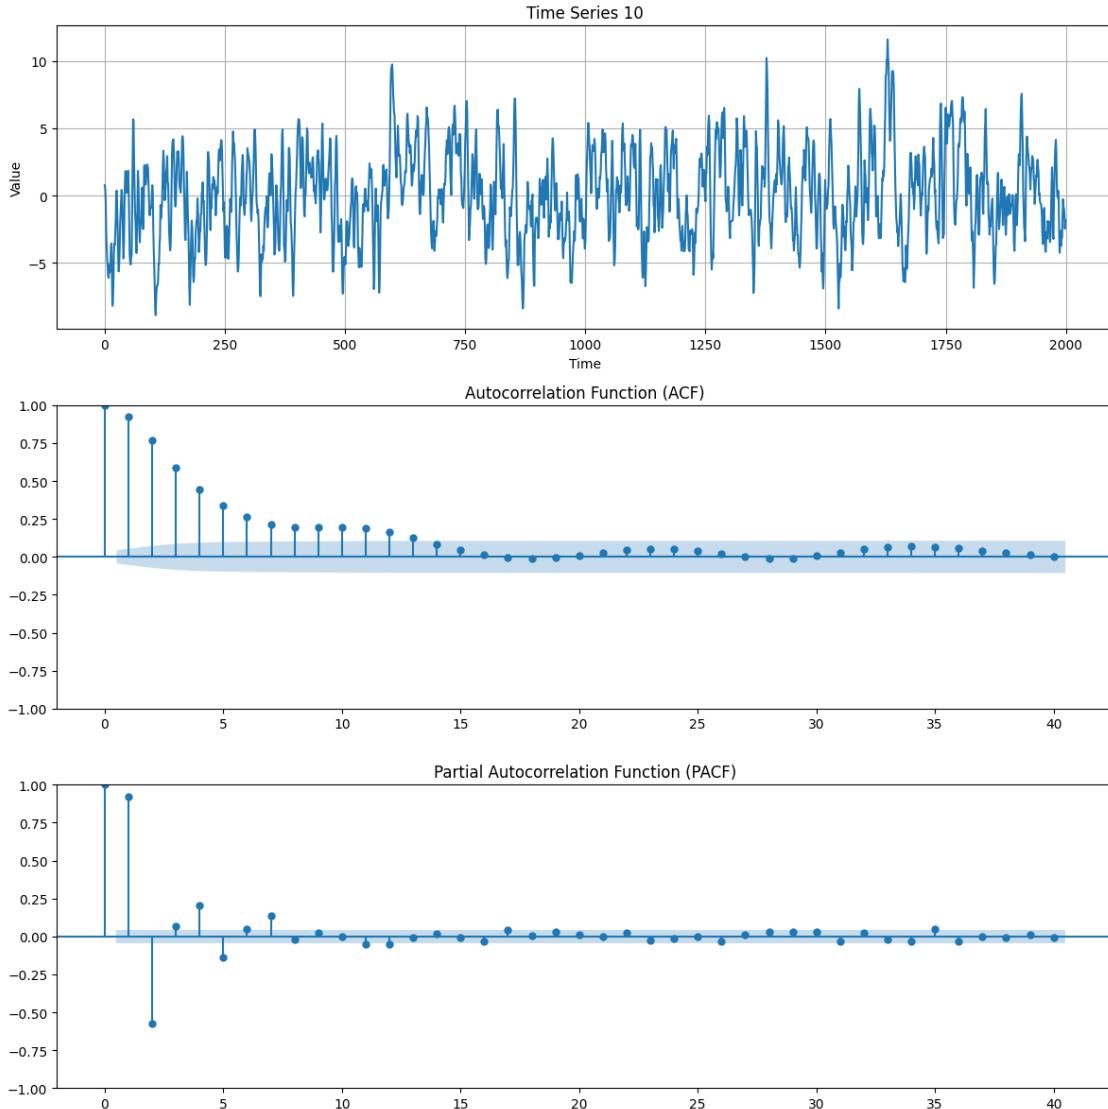
# Initialize the figure with 3 subplots
fig, (ax1, ax2, ax3) = plt.subplots(3, 1, figsize=(12, 12))

# Plot the Time Series
ax1.plot(df["t"], y)
ax1.set_title("Time Series 10")
ax1.set_xlabel("Time")
ax1.set_ylabel("Value")
ax1.grid(True)

# Plot Autocorrelation (ACF)
plot_acf(y, ax=ax2, lags=40, title="Autocorrelation Function (ACF)")

# Plot Partial Autocorrelation (PACF)
plot_pacf(y, ax=ax3, lags=40, title="Partial Autocorrelation Function (PACF)")

plt.tight_layout()
plt.show()
```



The autocorrelation structure of the series was analyzed through the Autocorrelation Function (ACF) and the Partial Autocorrelation Function (PACF).

The ACF exhibits a smooth and gradual decay, remaining significantly different from zero for several initial lags. This behavior is not compatible with a pure MA process, which would display a sharp cutoff in the ACF, and instead suggests the presence of autoregressive dynamics.

The PACF shows a small number of significant spikes at the first lags, followed by values that mostly lie within the confidence bounds. This pattern indicates that the dependence on past observations is captured by a low-order AR component.

Overall, the absence of a clear cutoff in either the ACF or the PACF suggests that neither a pure AR nor a pure MA model is sufficient. Instead, the observed patterns are consistent with a mixed **ARMA** structure with low orders.

Based on this preliminary identification step, ARMA models with small values of  $p$  and  $q$  are natural candidates for the series.

### 0.3 ARMA model identification

```
[4]: from statsmodels.tsa.arima.model import ARIMA

# Candidate ARMA models
candidates = [
    (1, 0, 0), # AR(1)
    (2, 0, 0), # AR(2)
    (1, 0, 1), # ARMA(1,1)
    (2, 0, 1), # ARMA(2,1)
    (1, 0, 2), # ARMA(1,2)
]

results = []

for order in candidates:
    try:
        # we remove the trend component by setting trend="n"
        model = ARIMA(y, order=order, trend="n")
        fit = model.fit()
        print(fit.summary())

        # we only have ARMA models here, we can ignore the 'd' parameter
        # so we use order[0] and order[2]
        results.append(
            {"model": f"ARMA({order[0]},{order[2]})", "AIC": fit.aic, "BIC": fit.bic}
        )
    except Exception as e:
        print(f"Model ARMA{order} failed: {e}")

results_df = pd.DataFrame(results).sort_values("AIC")
```

```
SARIMAX Results
=====
Dep. Variable:          series_10    No. Observations:             2000
Model:                  ARIMA(1, 0, 0)    Log Likelihood      -3299.519
Date: Mon, 15 Dec 2025   AIC                 6603.039
Time: 14:50:43           BIC                 6614.241
Sample:                0 - 2000           HQIC                6607.152
Covariance Type:         opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.9221	0.009	102.156	0.000	0.904	0.940
sigma2	1.5852	0.052	30.473	0.000	1.483	1.687
====						
====						
Ljung-Box (L1) (Q):			559.62	Jarque-Bera (JB):		
2.47						
Prob(Q):			0.00	Prob(JB):		
0.29						
Heteroskedasticity (H):			1.05	Skew:		
0.05						
Prob(H) (two-sided):			0.56	Kurtosis:		
2.86						
====						
====						
====						

## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## SARIMAX Results

Dep. Variable:	series_10	No. Observations:	2000
Model:	ARIMA(2, 0, 0)	Log Likelihood	-2902.154
Date:	Mon, 15 Dec 2025	AIC	5810.308
Time:	14:50:44	BIC	5827.110
Sample:	0 - 2000	HQIC	5816.477

Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	1.4503	0.018	78.947	0.000	1.414	1.486
ar.L2	-0.5724	0.018	-31.423	0.000	-0.608	-0.537
sigma2	1.0649	0.035	30.592	0.000	0.997	1.133
====						
====						
Ljung-Box (L1) (Q):			3.07	Jarque-Bera (JB):		
2.03						
Prob(Q):			0.08	Prob(JB):		
0.36						
Heteroskedasticity (H):			1.12	Skew:		
0.05						
Prob(H) (two-sided):			0.13	Kurtosis:		
2.88						
====						
====						

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

SARIMAX Results

```
=====
Dep. Variable:          series_10    No. Observations:             2000
Model:                 ARIMA(1, 0, 1)    Log Likelihood        -3021.391
Date:                 Mon, 15 Dec 2025   AIC                  6048.781
Time:                     14:50:44     BIC                  6065.584
Sample:                   0      HQIC                  6054.951
                           - 2000
Covariance Type:            opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.8802	0.011	76.665	0.000	0.858	0.903
ma.L1	0.4483	0.021	21.796	0.000	0.408	0.489
sigma2	1.2000	0.040	30.238	0.000	1.122	1.278

<=====

==

Ljung-Box (L1) (Q): 46.02 Jarque-Bera (JB):

4.16

Prob(Q): 0.00 Prob(JB):

0.12

Heteroskedasticity (H): 1.12 Skew:

0.07

Prob(H) (two-sided): 0.14 Kurtosis:

2.83

=====

==

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

SARIMAX Results

```
=====
Dep. Variable:          series_10    No. Observations:             2000
Model:                 ARIMA(2, 0, 1)    Log Likelihood        -2899.557
Date:                 Mon, 15 Dec 2025   AIC                  5807.115
Time:                     14:50:44     BIC                  5829.518
Sample:                   0      HQIC                  5815.341
                           - 2000
Covariance Type:            opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	1.4063	0.032	43.426	0.000	1.343	1.470
ar.L2	-0.5319	0.031	-17.260	0.000	-0.592	-0.472

<=====

ma.L1	0.0666	0.037	1.777	0.076	-0.007	0.140
sigma2	1.0622	0.035	30.560	0.000	0.994	1.130

=====

Ljung-Box (L1) (Q):	2.08	0.15	Jarque-Bera (JB):
Prob(Q):	0.35	0.70	Prob(JB):
Heteroskedasticity (H):	0.05	1.13	Skew:
Prob(H) (two-sided):	2.87	0.11	Kurtosis:

=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

#### SARIMAX Results

Dep. Variable:	series_10	No. Observations:	2000
Model:	ARIMA(1, 0, 2)	Log Likelihood	-2822.638
Date:	Mon, 15 Dec 2025	AIC	5653.275
Time:	14:50:45	BIC	5675.679
Sample:	0 - 2000	HQIC	5661.502

Covariance Type:	opg
------------------	-----

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.7899	0.016	48.656	0.000	0.758	0.822
ma.L1	0.7050	0.021	33.190	0.000	0.663	0.747
ma.L2	0.5075	0.021	24.294	0.000	0.467	0.548
sigma2	0.9834	0.032	30.905	0.000	0.921	1.046

=====

Ljung-Box (L1) (Q):	1.54	0.73	Jarque-Bera (JB):
Prob(Q):	0.46	0.39	Prob(JB):
Heteroskedasticity (H):	0.05	1.14	Skew:
Prob(H) (two-sided):	2.91	0.10	Kurtosis:

=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

	model	AIC	BIC
4	ARMA(1,2)	5653.275490	5675.679100
3	ARMA(2,1)	5807.114706	5829.518315
1	ARMA(2,0)	5810.307591	5827.110298
2	ARMA(1,1)	6048.781362	6065.584069
0	ARMA(1,0)	6603.038797	6614.240602

Since the series is stationary, an ARMA model can be directly fitted. The identification follows the Box–Jenkins methodology.

The inspection of the ACF and PACF suggests the presence of both autoregressive and moving-average components with relatively low orders. Based on this preliminary analysis, several candidate ARMA models were estimated.

Model comparison was carried out using information criteria. The table above reports the AIC and BIC values for all candidate models. Among them, the **ARMA(1,2)** (without constant) model achieves the lowest values of both AIC and BIC, indicating the best balance between goodness of fit and model parsimony.

Furthermore, all parameters of the ARMA(1,2) model are statistically significant, and the Ljung–Box test applied to the residuals does not reject the null hypothesis of no autocorrelation. This indicates that the residuals behave like white noise.

Therefore, **ARMA(1,2)** is selected as the final model for the series.

### 0.3.1 Residual normality

The normality of the residuals was assessed using the Jarque–Bera test.

The test yields a p-value well above standard significance levels. Therefore, we do not reject the null hypothesis of normality. This provides no evidence against the assumption that the residuals are normally distributed.

Although normality is not strictly required for the consistency of ARMA estimators, this result supports the adequacy of the model and facilitates statistical inference.