

Group10 Second Assignment (part 2)

December 16, 2025

Authors

Mattia Zanin - mattia.zanin@studenti.unipd.it
Matteo Giorgi - matteo.giorgi.1@studenti.unipd.it
Enrico Zanello - enrico.zanello@studenti.unipd.it
Luca Lo Buono - luca.lobuono@studenti.unipd.it

0.1 Train/Test split (time-ordered)


We evaluate forecasting performance out-of-sample using a time-ordered split. The last H observations are kept as test set, while the remaining observations are used for training. This mirrors the real forecasting setting where future data are not available at estimation time.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

df = pd.read_csv("arma12_seasonal_controlled_15series.csv")
y = df["series_10"].astype(float).to_numpy()

H = 200 # test horizon (typical choice: 100-300 depending on length)
y_train, y_test = y[:-H], y[-H:]

plt.figure(figsize=(12, 4))
plt.plot(y_train, label="train")
plt.plot(np.arange(len(y_train), len(y)), y_test, label="test")
plt.title("Train/Test split - Group 10")
plt.grid(alpha=0.3)
plt.legend()
plt.show()
```



0.2 Model fitting

Following the Box–Jenkins identification performed in the previous assignment, we fit the selected ARMA(p,q) model on the training set.

```
[2]: from statsmodels.tsa.arima.model import ARIMA

# ARMA(1,2) = ARIMA(1, 0, 2)
fit = ARIMA(y_train, order=(1, 0, 2), trend="n").fit()
print(fit.summary())
```

| SARIMAX Results | | | | | | |
|---------------------|------------------|-------------------|-------------------|-------|--------|--------|
| ===== | | | | | | |
| Dep. Variable: | y | No. Observations: | 1800 | | | |
| Model: | ARIMA(1, 0, 2) | Log Likelihood | -2536.299 | | | |
| Date: | Tue, 16 Dec 2025 | AIC | 5080.598 | | | |
| Time: | 10:28:11 | BIC | 5102.580 | | | |
| Sample: | 0 - 1800 | HQIC | 5088.713 | | | |
| Covariance Type: | opg | | | | | |
| ===== | | | | | | |
| | coef | std err | z | P> z | [0.025 | 0.975] |
| ===== | | | | | | |
| ar.L1 | 0.7972 | 0.017 | 47.577 | 0.000 | 0.764 | 0.830 |
| ma.L1 | 0.7062 | 0.022 | 31.449 | 0.000 | 0.662 | 0.750 |
| ma.L2 | 0.5113 | 0.022 | 23.595 | 0.000 | 0.469 | 0.554 |
| sigma2 | 0.9787 | 0.034 | 29.064 | 0.000 | 0.913 | 1.045 |
| ===== | | | | | | |
| Ljung-Box (L1) (Q): | 1.57 | 0.67 | Jarque-Bera (JB): | | | |
| Prob(Q): | 0.46 | 0.41 | Prob(JB): | | | |

```

Heteroskedasticity (H):           1.09      Skew:
0.04
Prob(H) (two-sided):            0.30      Kurtosis:
2.89
=====
=====
```

====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

0.3 Point forecasts and accuracy

We produce H-step-ahead forecasts for the test period and evaluate accuracy using standard loss functions (MAE and RMSE).


```
[3]: pred = fit.get_forecast(steps=H)
y_hat = pred.predicted_mean

mae = np.mean(np.abs(y_test - y_hat))
rmse = np.sqrt(np.mean((y_test - y_hat) ** 2))
print("MAE:", mae)
print("RMSE:", rmse)

plt.figure(figsize=(12, 4))
plt.plot(y_test, label="test")
plt.plot(y_hat, label="forecast")
plt.title("Out-of-sample forecasts (point)")
plt.grid(alpha=0.3)
plt.legend()
plt.show()
```

MAE: 2.10458178927103

RMSE: 2.637130596618867



0.3.1 Rolling 1-step (se richiesto / consigliato)

```
[4]: from statsmodels.tsa.arima.model import ARIMA

history = list(y_train)
roll_forecasts = []

for t in range(H):
    m = ARIMA(history, order=(1, 0, 2), trend="n").fit()
    roll_forecasts.append(m.forecast(steps=1)[0])
    history.append(y_test[t])

roll_forecasts = np.array(roll_forecasts)

mae_r = np.mean(np.abs(y_test - roll_forecasts))
rmse_r = np.sqrt(np.mean((y_test - roll_forecasts) ** 2))
print("Rolling MAE:", mae_r)
print("Rolling RMSE:", rmse_r)
```

Rolling MAE: 0.8045793501230344

Rolling RMSE: 1.0131027773912855


0.4 Interval forecasts

We compute 95% prediction intervals. In a correct ARMA specification, most test observations should fall within these bands, and the interval width reflects forecast uncertainty increasing with the horizon.

```
[5]: ci = pred.conf_int(alpha=0.05) # 95%
lower = ci[:, 0]
upper = ci[:, 1]

plt.figure(figsize=(12, 4))
plt.plot(y_test, label="test")
plt.plot(y_hat, label="forecast")
plt.fill_between(np.arange(H), lower, upper, alpha=0.2, label="95% PI")
plt.title("Forecast with 95% prediction intervals")
plt.grid(alpha=0.3)
plt.legend()
plt.show()

coverage = np.mean((y_test >= lower) & (y_test <= upper))
print("Empirical coverage:", coverage)
```



Empirical coverage: 0.98

0.5 Benchmark comparison

We compare ARMA forecasts against a simple naive benchmark: $\hat{y}_{t+1} = y_t$. Despite its simplicity, the naive model is often hard to beat in short-horizon forecasting.

```
[6]: naive = np.r_[y_train[-1], y_test[:-1]] # forecast at t uses previous observed value
mae_n = np.mean(np.abs(y_test - naive))
rmse_n = np.sqrt(np.mean((y_test - naive) ** 2))
print("Naive MAE:", mae_n)
print("Naive RMSE:", rmse_n)
```

Naive MAE: 1.015481502653484
 Naive RMSE: 1.2632077429018165